

Clasificación de imágenes con TensorFlow

...

Alejandro Solano - Málaga Python



MÁLAGA
PYTHON







gato



input



target

??

gato



input



target

sin
+
X
log
exp

gato




input

target



\sin
 $+$
 \times
 \log
 \exp

gato



TensorFlow

Deep Learning





MODEL




gato



Red neuronal


$$\mathbf{x} \cdot \mathbf{W}_1 + \mathbf{b}_1 = \mathbf{y}$$



Red neuronal

$$x \cdot w_1 + b_1 = y$$

$$(x \cdot w_1 + b_1) \cdot w_2 + b_2 = y$$




Red neuronal

$$\mathbf{x} \cdot \mathbf{w}_1 + \mathbf{b}_1 = y$$

$$(\mathbf{x} \cdot \mathbf{w}_1 + \mathbf{b}_1) \cdot \mathbf{w}_2 + \mathbf{b}_2 = y$$

$$((\mathbf{x} \cdot \mathbf{w}_1 + \mathbf{b}_1) \cdot \mathbf{w}_2 + \mathbf{b}_2) \cdot \mathbf{w}_3 + \mathbf{b}_3 = y$$




Red neuronal

$$x \cdot W_1 + b_1 = y$$


$$(x \cdot W_1 + b_1) \cdot W_2 + b_2 = y$$

$$\tanh(\sigma(x \cdot W_1 + b_1) \cdot W_2 + b_2) \cdot W_3 + b_3 = y$$



Redes neuronales: intuición

Las primeras capas extraen información más básica, y las siguientes trabajan a partir de esa información.



hace la suma


clasifica la suma




Redes neuronales: intuición

feature
extractors

classifiers



Redes neuronales: intuición



Redes neuronales: intuición

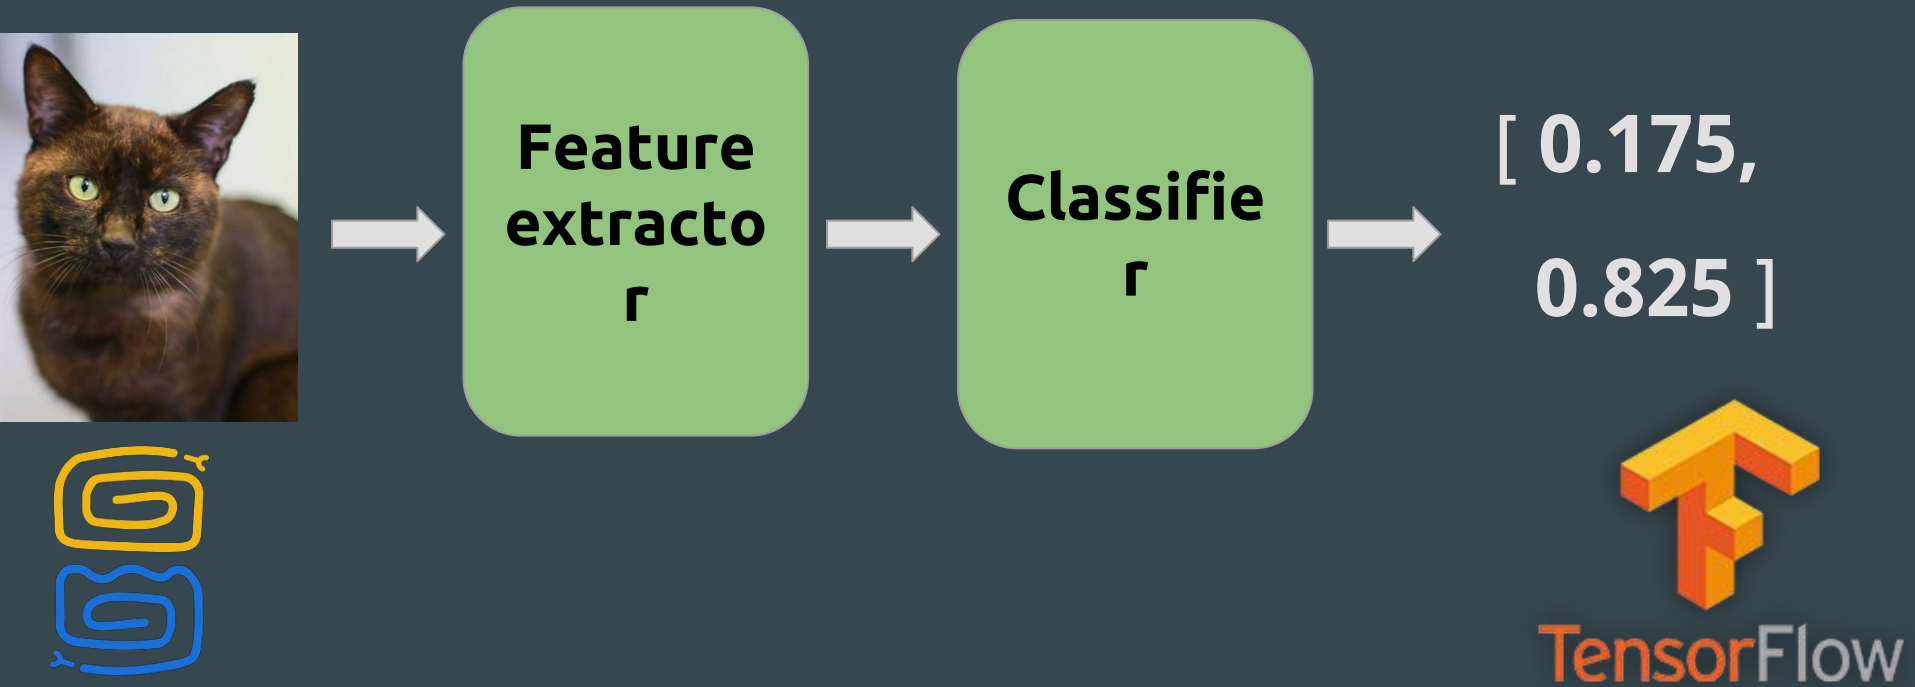


Image Classification






Features: intuition

¿Qué hace que un gato sea un gato?






Features: intuition


¿Qué hace que un gato sea un gato?



Features: intuition



**Classifie
r**





gato



TensorFlow

Features: intuition



**Classifie
r**




caballo







TensorFlow

Features: intuition




**Feature
extracto
r**



TensorFlow

Analyzing input data: Images




Conjunto de píxeles

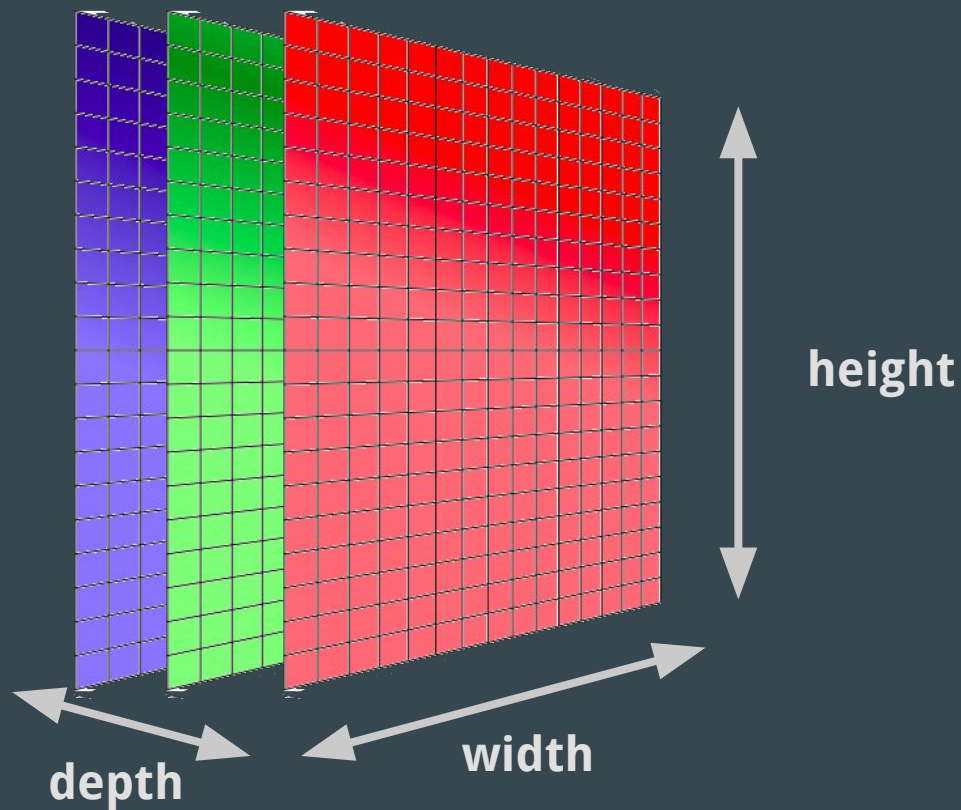
**Cada píxel toma
valores de 0 a 255**




Analyzing input data: Images



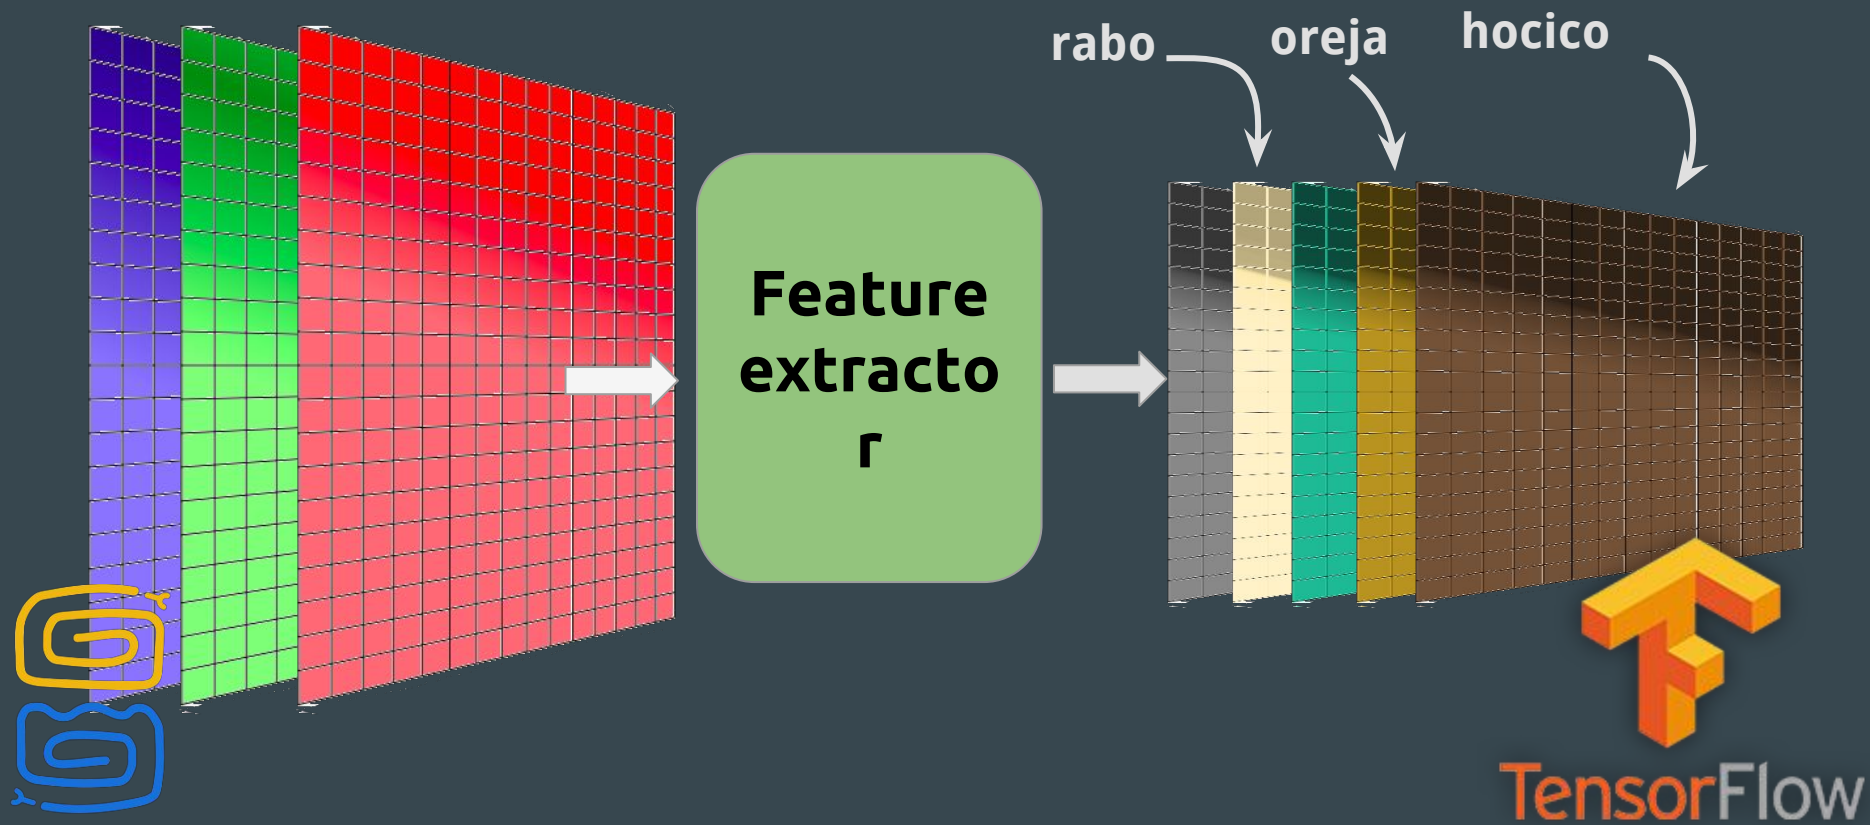
Analyzing input data: Images




Analyzing input data: Images




Analyzing input data: Images




Analyzing input data: Images




Analyzing input data: Images



Analyzing input data: Images



Analyzing input data: Images



Analyzing input data: Images



TensorFlow

Analyzing input data: Images



TensorFlow

Analyzing input data: Images




TensorFlow

Analyzing input data: Images




TensorFlow


Convolution



Convolution




Convolution



Convolution


kernel



TensorFlow

Convolution

kernel




**each kernel learns to
extract one feature**




TensorFlow

Convolution



Convolution




Convolution




TensorFlow

Strides and Padding




Strides and Padding

stride 1



Strides and Padding


stride 1



Strides and Padding

stride 1


stride 2




Strides and Padding

stride 1

stride 2



Strides and Padding




same padding



TensorFlow

Strides and Padding



same padding

valid padding



Convolutional layer: code

```
# kernels
W = tf.Variable(tf.truncated_normal(
    [kernel_size, kernel_size, input_depth, num_kernels],
    stddev=0.05))

b = tf.Variable(tf.zeros([num_kernels]))
```




Convolutional layer: code

```
# convolutional layer
conv = tf.nn.conv2d(x, W,
                    strides=[1, strides, strides, 1],
                    padding='SAME')

conv = tf.nn.bias_add(conv, b)
```



Non-linear function: ReLu



TensorFlow

Non-linear function: ReLu

```
conv = tf.nn.relu(conv)
```



Max-Pooling




TensorFlow

Max-Pooling



TensorFlow

Max-Pooling



TensorFlow

Max-Pooling



Reduzco el 75% de los parámetros

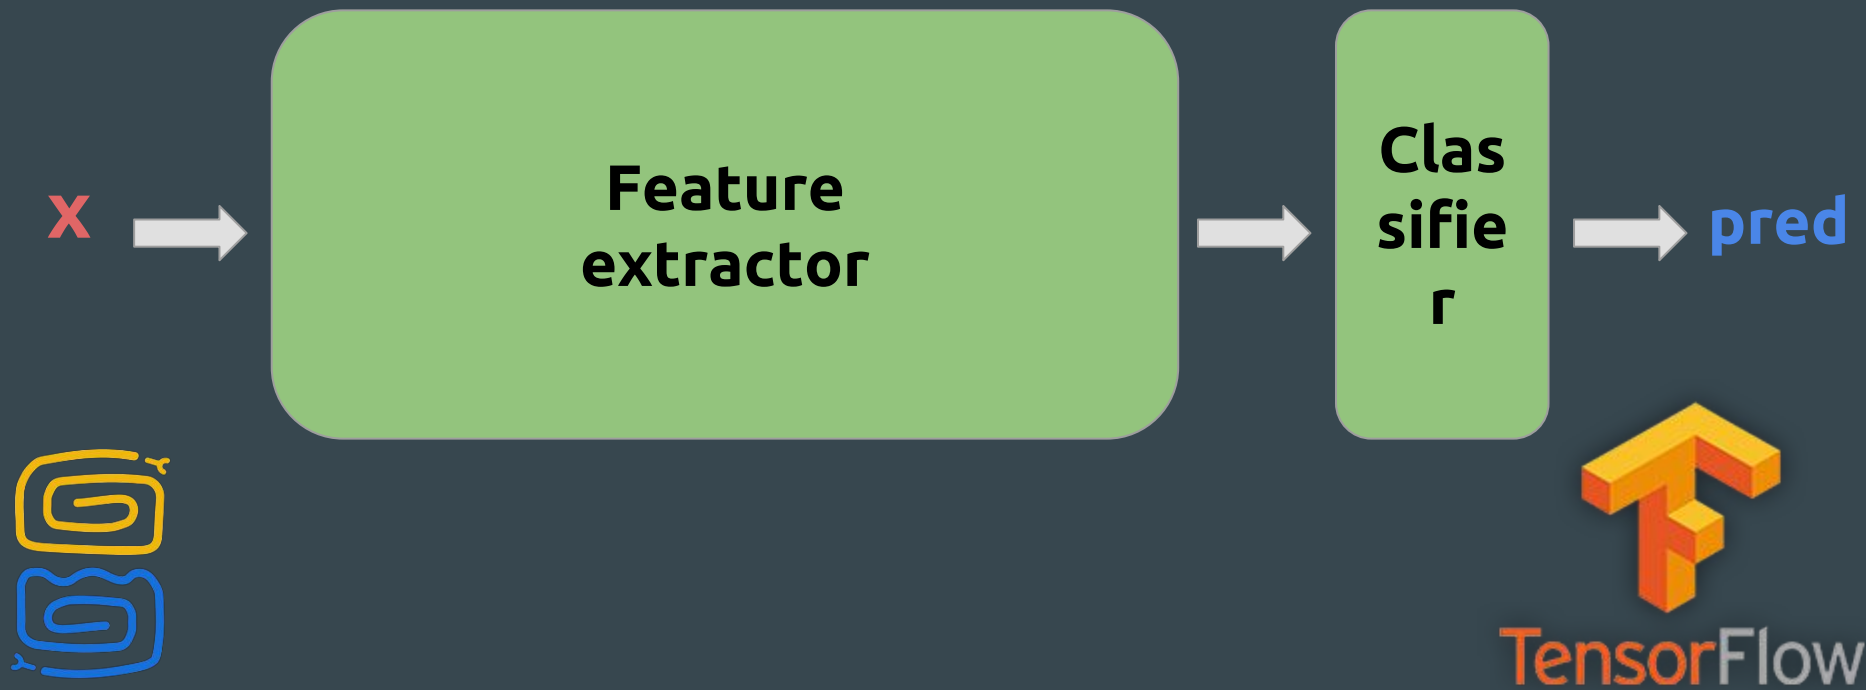


Max-Pooling


```
x = tf.nn.max_pool(x,  
    ksize=[1, k, k, 1],  
    strides=[1, k, k, 1],  
    padding='SAME')
```




Flatten Layer




Flatten Layer



Flatten Layer



Flatten Layer



Flatten Layer

```
x_shape = x.get_shape().as_list()
```

```
flatten_shape = x_shape[1] * x_shape[2] * x_shape[3]
```

```
flatten = tf.reshape(x, [-1, flatten_shape])
```



CIFAR-10



CIFAR-10

plane



car




bird



cat



deer



dog



frog



horse



ship



truck



CIFAR-10



32x32x3



**Feature
extractor**

**Clas
sifi
er**

0.1
0
0.05
0.1
0
0
0.1
0.6
0
0.05

CIFAR-10

```
# PLACEHOLDERS  
x = tf.placeholder(tf.float32, [None, 32, 32, 3])  
y = tf.placeholder(tf.float32, [None, 10])
```



Preprocessing

input



standardize




$-\mu$

σ

target

7

one-hot encode



$(0, 0, 0, 0, 0, 0, 0, 1, 0, 0)$

Cost Function: cross-entropy

$$C = -\frac{1}{n} \sum_x [y \ln a + (1 - y) \ln(1 - a)],$$



típica en clasificación




Cost Function: cross-entropy

```
cost =  
tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(  
logits=logits, labels=y))
```



Adam Optimizer

ajusta el learning rate



TensorFlow

Adam Optimizer

```
optimizer = tf.train.AdamOptimizer().minimize(cost)
```



Vamos al notebook



¿Es mejorable ese 70%?




Data augmentation



original



flip



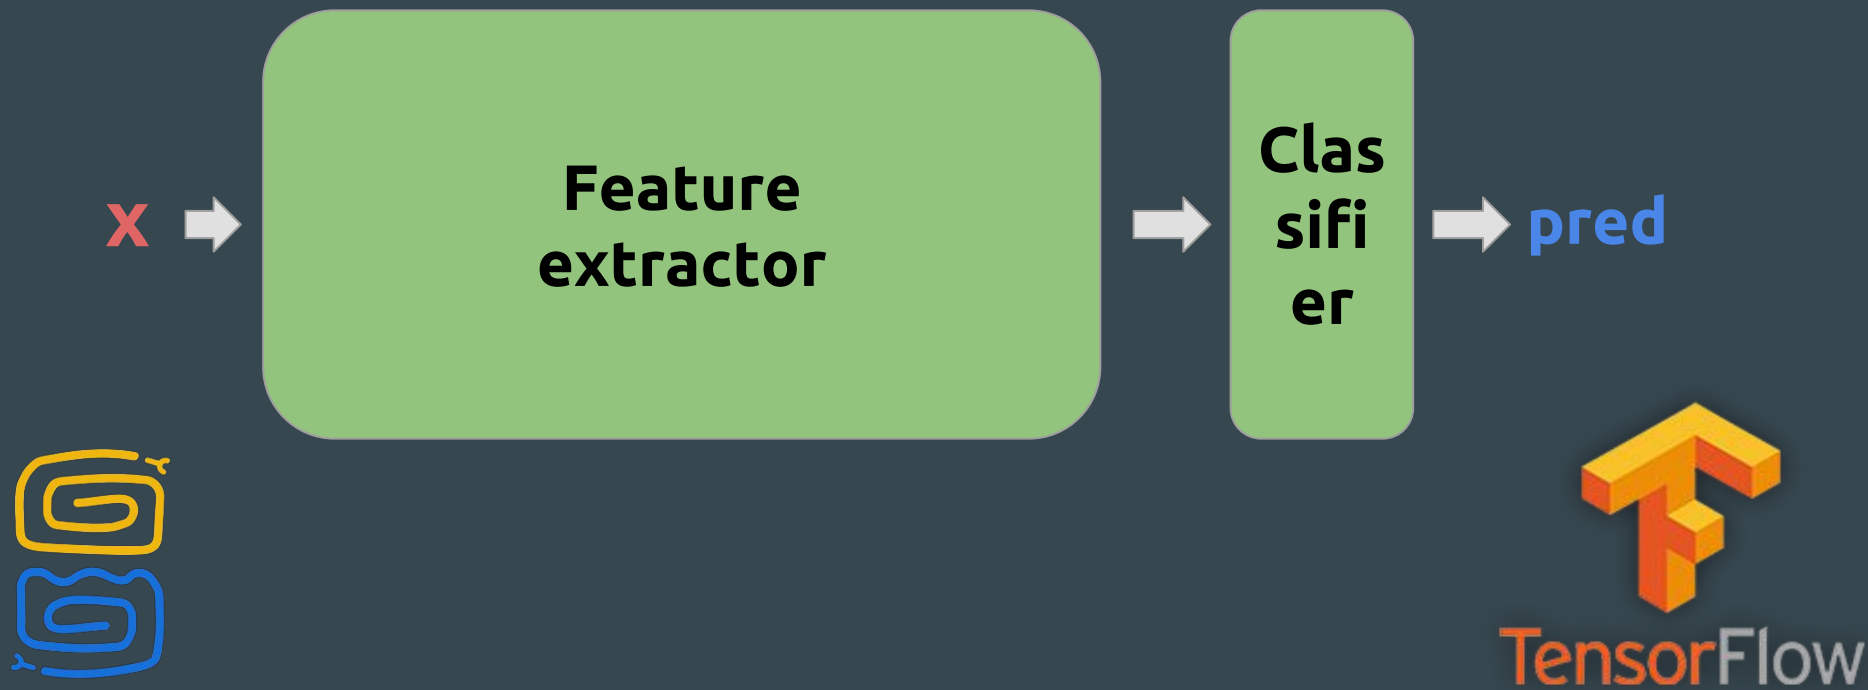
zoom



brighten



Enlarging the network



Enlarging the network



Enlarging the network

x →

**Feature
extractor**

→

**Cl
si
e**



TensorFlow


Transfer Learning



ImageNet



VGG-16

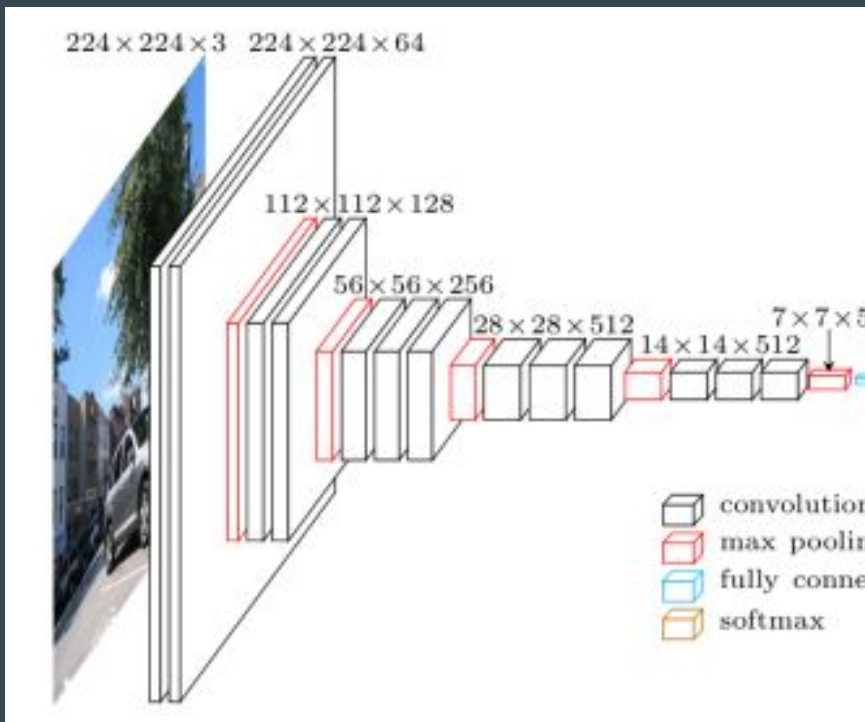


TensorFlow

transfer learning



x



Classifier



pred




TensorFlow

transfer learning



x




x'



TensorFlow

transfer learning



Vamos al notebook



Gracias!





alesolano/mastering_tensorflow

